

Backpropagation

手推bp的流程

Deriving gradients: Tips

- **Tip 1:** Carefully define your variables and keep track of their dimensionality!
- **Tip 2:** Chain rule! If $\mathbf{y} = f(\mathbf{u})$ and $\mathbf{u} = g(\mathbf{x})$, i.e., $\mathbf{y} = f(g(\mathbf{x}))$, then:

$$\frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \frac{\partial \mathbf{y}}{\partial \mathbf{u}} \frac{\partial \mathbf{u}}{\partial \mathbf{x}}$$

Keep straight what variables feed into what computations

- **Tip 3:** For the top softmax part of a model: First consider the derivative wrt f_c when $c = y$ (the correct class), then consider derivative wrt f_c when $c \neq y$ (all the incorrect classes)
- **Tip 4:** Work out element-wise partial derivatives if you're getting confused by matrix calculus!
- **Tip 5:** Use Shape Convention. Note: The error message δ that arrives at a hidden layer has the same dimensionality as that hidden layer

6

下游任务是否用word vec输入

So what should I do?

- **Question:** Should I use available “pre-trained” word vectors
Answer:
 - Almost always, yes!
 - They are trained on a huge amount of data, and so they will know about words not in your training data and will know more about words that are in your training data
 - Have 100s of millions of words of data? Okay to start random
- **Question:** Should I update (“fine tune”) my own word vectors?
- **Answer:**
 - If you only have a **small** training data set, **don't** train the word vectors
 - If you have have a **large** dataset, it probably will work better to **train = update = fine-tune** word vectors to the task

11

上图说明了wordvec是否应该在下游任务中更新。

Computation Graphs and Backpropagation

用图表示神经网络结构，下图是forward propagation

Computation Graphs and Backpropagation

- We represent our neural net equations as a graph

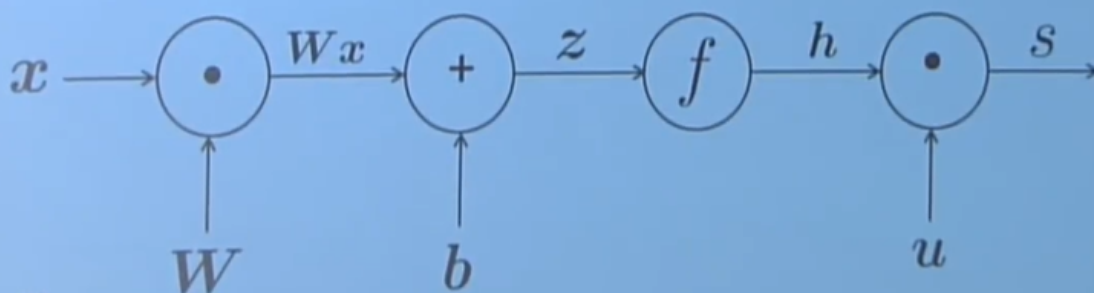
- Source nodes: inputs
- Interior nodes: operations
- Edges pass along result of the operation

$$s = u^T h$$

$$h = f(z)$$

$$z = Wx + b$$

$$x \text{ (input)}$$



14

下图是backpropagation

Backpropagation

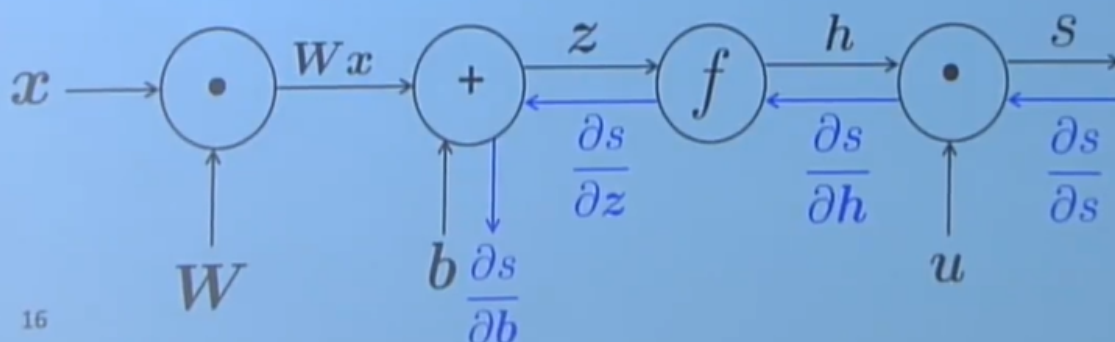
- Go backwards along edges
 - Pass along **gradients**

$$s = u^T h$$

$$h = f(z)$$

$$z = Wx + b$$

$$x \text{ (input)}$$



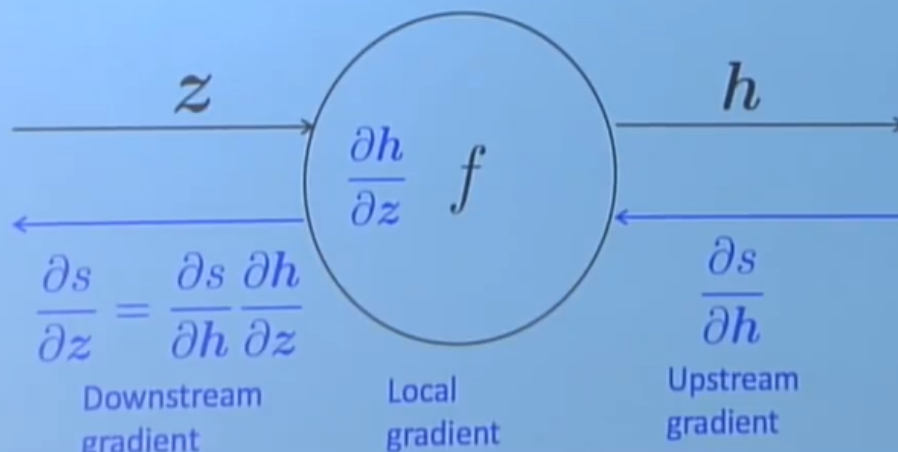
16

下图是用chain rule计算多节点的梯度，重要的公式 $downstream = upstream * local$

Backpropagation: Single Node

- Each node has a **local gradient**
 - The gradient of it's output with respect to it's input
- [downstream gradient] = [upstream gradient] x [local gradient]

$$h = f(z)$$

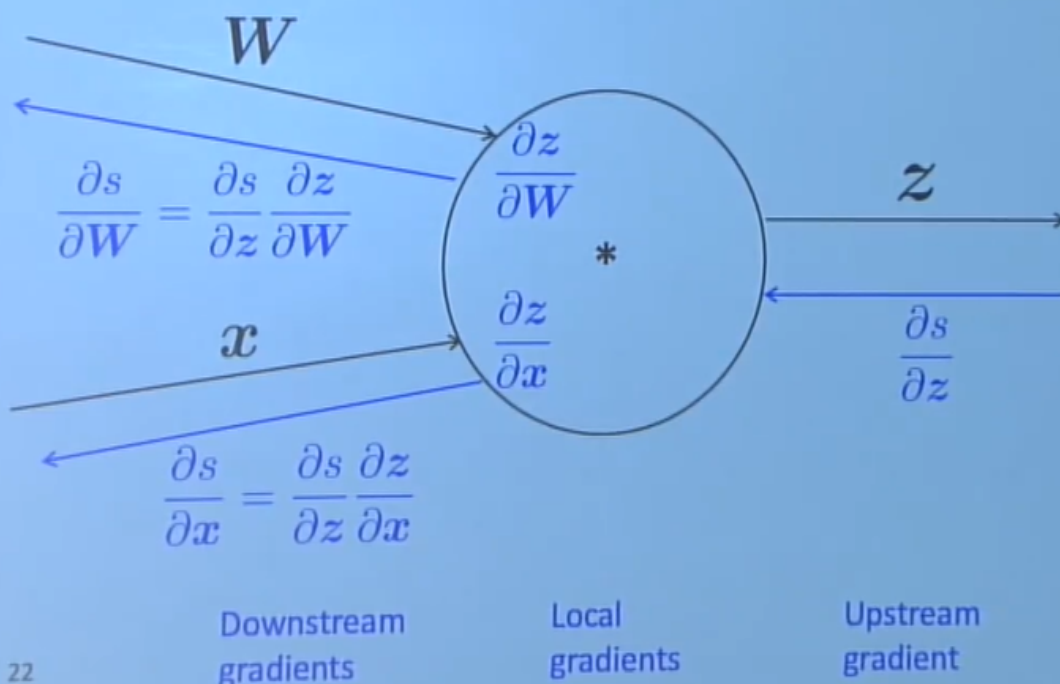


20

Backpropagation: Single Node

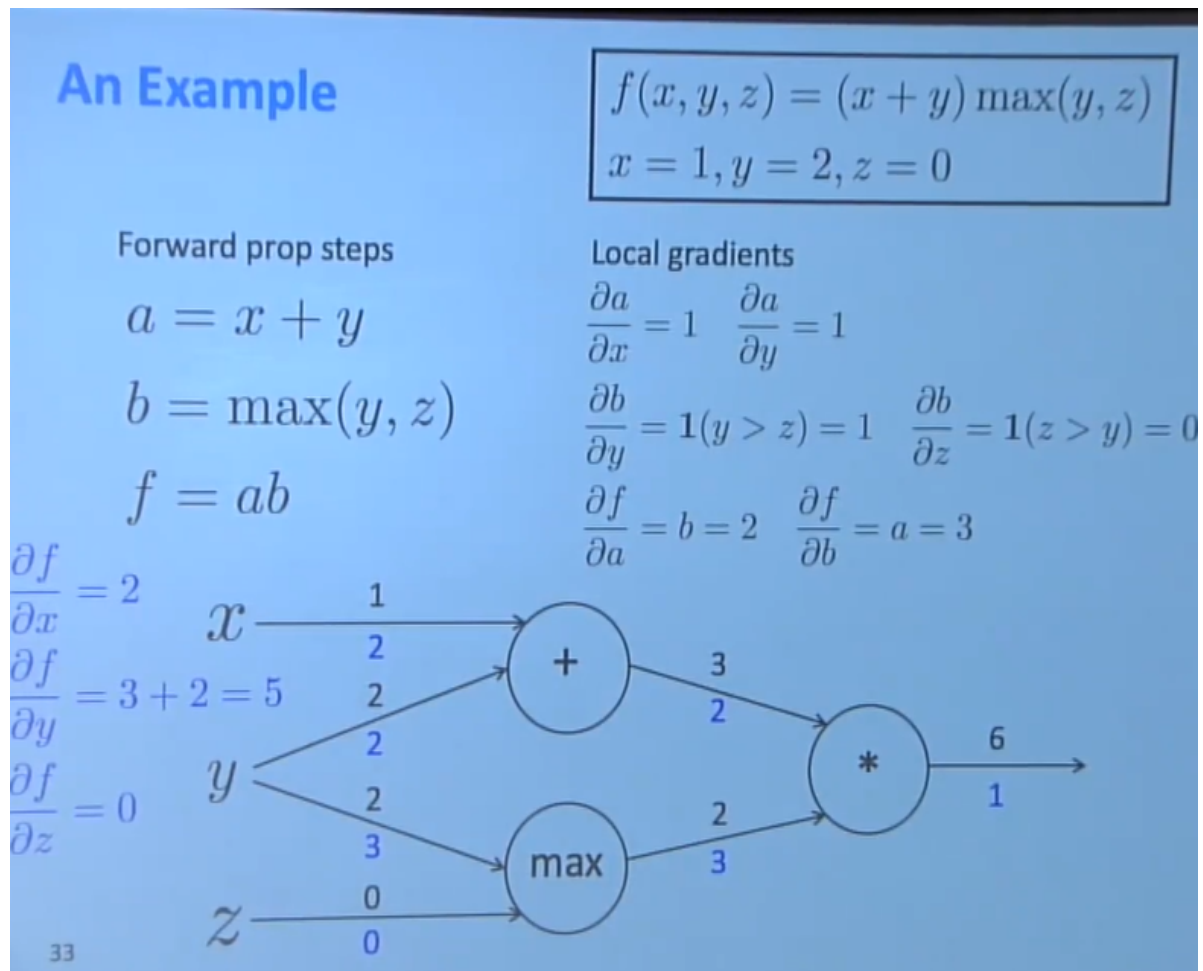
- Multiple inputs \rightarrow multiple local gradients

$$z = Wx$$



22

一个backpropagation例子



能算出左边的结果即可， $\frac{\partial f}{\partial x}$ 物理意义是x改变0.1，f就会改变0.2。

Efficiency: compute all gradients at once

- Correct way:

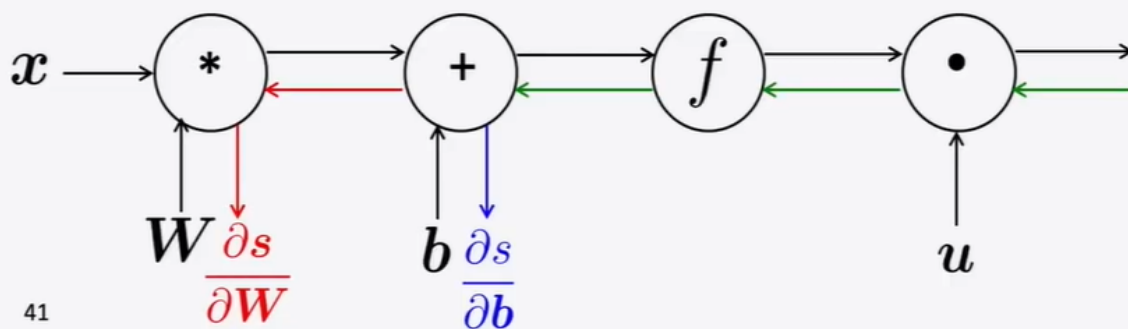
- Compute all the gradients at once
- Analogous to using δ when we computed gradients by hand

$$s = u^T h$$

$$h = f(z)$$

$$z = Wx + b$$

$$x \quad (\text{input})$$



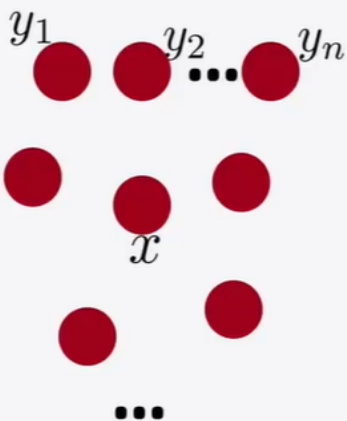
41

上图说明了，上游的梯度(绿色)可以保留，不用重复计算。

Back-Prop in General Computation Graph

Single scalar output z

...



1. Fprop: visit nodes in topological sort order
 - Compute value of node given predecessors
2. Bprop:
 - initialize output gradient = 1
 - visit nodes in reverse order:

Compute gradient wrt each node using gradient wrt successors

$\{y_1, y_2, \dots, y_n\} = \text{successors of } x$

$$\frac{\partial z}{\partial x} = \sum_{i=1}^n \frac{\partial z}{\partial y_i} \frac{\partial y_i}{\partial x}$$

Done correctly, big $O()$ complexity of fprop and bprop is **the same**

In general our nets have regular layer-structure and so we can use matrices and Jacobians...

42

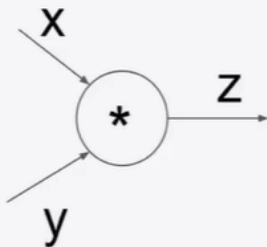
上图是bp的构图过程和复杂度。

Backprop Implementations

```
class ComputationalGraph(object):
    #...
    def forward(inputs):
        # 1. [pass inputs to input gates...]
        # 2. forward the computational graph:
        for gate in self.graph.nodes_topologically_sorted():
            gate.forward()
        return loss # the final gate in the graph outputs the loss
    def backward():
        for gate in reversed(self.graph.nodes_topologically_sorted()):
            gate.backward() # little piece of backprop (chain rule applied)
        return inputs_gradients
```

44

Implementation: forward/backward API



(x,y,z are scalars)

```
class MultiplyGate(object):
    def forward(x,y):
        z = x*y
        self.x = x # must keep these around!
        self.y = y
        return z
    def backward(dz):
        dx = self.y * dz # [dz/dx * dL/dz]
        dy = self.x * dz # [dz/dy * dL/dz]
        return [dx, dy]
```

上面两张图是深度学习框架图计算的流程。